

TELE6420 INFRASTRUCTURE AUTOMATION DESIGN AND TOOLS

FINAL PROJECT

Project By:

Shah Nawaz Syed Shah

Ali Goudarzi

Rohit Wagh

CONTENTS:

1. CREDENTIAL AND LINKS
2. CODE EXPLANATION WITH SCREENSHOTS
 - Region & AZ
 - VPC
 - Internet gateway and route table
 - Security group "WEB_SG"
 - Launch configuration
 - Application Load Balancer
 - Route 53
 - Auto Scaling Group
3. CPU LOAD
4. WEB APPLICATION
5. DOMAIN NAME
6. DIFFICULTIES
7. BONUS
8. FUTURE SCOPE

CREDENTIAL AND LINKS:

Website: <http://projiadt.online>

AWS login link: <https://996278926886.signin.aws.amazon.com/console>

Iam_user name: proj_prof

Password: Hihello123.

Access Key: AKIA6P5WZ7YTLNXIC7NZ

Shared Access Key: b0tNwe4SooiLswsulxPPe44ZZWKvyVXxKMkeFU9D

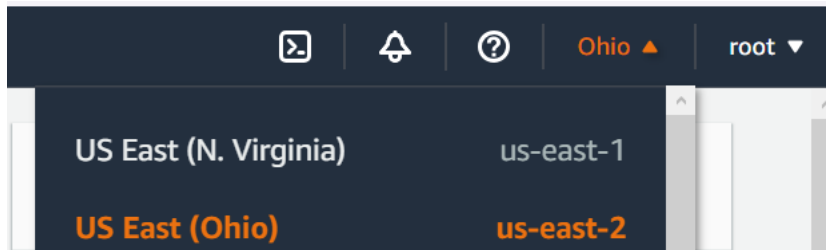
Github: <https://github.com/finalprojiadt/iadt>

Dockerhub: <https://hub.docker.com/repository/docker/projiadt/weeb/general>

CODE EXPLANATION:

Region & AZ

```
provider "aws" {  
  region = "us-east-2"  
}  
  
data "aws_availability_zones" "available" {  
  state = "available"  
}
```



The AWS provider is set up with a specified region of us-east-2. Additionally, to ensure high availability, all the Availability Zones (Az) within that region is used for data retrieval.

VPC

```
# Create VPC  
resource "aws_vpc" "main" {  
  cidr_block      = "10.0.0.0/16"  
  enable_dns_hostnames = true  
  enable_dns_support   = true  
  
  tags = {  
    Name = "Main VPC"  
  }  
}  
  
# Create 3 subnets in the VPC  
resource "aws_subnet" "subnets" {  
  count              = 3  
  vpc_id             = aws_vpc.main.id  
  cidr_block         = "10.0.${count.index}.0/24"  
  availability_zone   = element(data.aws_availability_zones.available.names,  
count.index)  
  map_public_ip_on_launch = count.index < 2 ? true : false  
  
  tags = {  
    Name = "Subnet-${count.index}"  
  }  
}
```

A new VPC named "Main VPC" has been established with a CIDR block of 10.0.0.0/16. This VPC is further segmented into three subnets with CIDR blocks: 10.0.0.0/24, 10.0.1.0/24, and 10.0.2.0/24, distributed across all available zones in us-east-2. The first two subnets are designated as public, meaning they automatically assign public IPs to instances, while the third subnet is private.

vpc-0ee45b8b41c5fbeb0b / Main VPC

Details Info

VPC ID vpc-0ee45b8b41c5fbeb0b	State Available	DNS hostnames Enabled	DNS resolution Enabled
Tenancy Default	DHCP option set dopt-06fe7bc83b4d0472e	Main route table rtb-06a82450b99339ed8	Main network ACL acl-0cb0589637e0c1a98
Default VPC No	IPv4 CIDR 10.0.0.0/16	IPv6 pool -	IPv6 CIDR -
Network Address Usage metrics Disabled	Route 53 Resolver DNS Firewall rule groups -	Owner ID 996278926886	

Resource map New CIDRs Flow logs Tags

Resource map Info

VPC Show details
Your AWS virtual network

Main VPC

Subnets (3)
Subnets within this VPC

us-east-2a
Subnet-0

us-east-2b
Subnet-1

us-east-2c
Subnet-2

Route tables (3)
Route network traffic to resources

rtb-06a82450b99339ed8

Private Route Table

Public Route Table

Network connections (1)
Connections to other networks

igw

Introducing the VPC resource map
Solid lines represent relationships between resources in your VPC. Dotted lines represent network traffic to external destinations.

Subnets (1/6) Info

Find resources by attribute or tag

Name	Subnet ID	State	VPC	IPv4 CIDR	IPv6 CIDR	Available IPv4
Subnet-0	subnet-0597c15362ddcb84e	Available	vpc-0ee45b8b41c5fbeb0b Main...	10.0.0.0/24	-	249
Subnet-2	subnet-02dc9a55261263444	Available	vpc-0ee45b8b41c5fbeb0b Main...	10.0.2.0/24	-	251
Subnet-1	subnet-0f7a149ea3434e50f	Available	vpc-0ee45b8b41c5fbeb0b Main...	10.0.1.0/24	-	250

Internet gateway and route table

```
# Create Internet Gateway and attach to VPC
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main.id

  tags = {
    Name = "igw"
  }
}

# Create route table and add public route
resource "aws_route_table" "public" {
  vpc_id = aws_vpc.main.id

  route {
```

```

    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.igw.id
  }
  tags = {
    Name = "Public Route Table"
  }
}
resource "aws_route_table" "private" {
  vpc_id = aws_vpc.main.id
  tags = {
    Name = "Private Route Table"
  }
}
resource "aws_route_table_association" "public" {
  count = 2

  subnet_id      = element(aws_subnet.subnets[*].id, count.index)
  route_table_id = aws_route_table.public.id
}
resource "aws_route_table_association" "private" {
  subnet_id      = aws_subnet.subnets[2].id
  route_table_id = aws_route_table.private.id
}

```

To enable instances to access the internet, an internet gateway and a public route table have been set up and associated with the first two subnets. Additionally, there's a private route table designed specifically for the private subnet, ensuring it's intended for internal network access only.

VPC > Internet gateways > igw-0236104df2b447193

igw-0236104df2b447193 / igw Actions ▼

Details [Info](#)

Internet gateway ID igw-0236104df2b447193	State Attached	VPC ID vpc-0ee45b8b41c5f8e0b Main VPC	Owner 996278926886
--	--------------------------------	--	-----------------------

Tags Manage tags

Search tags

Key	Value
Name	igw

Route tables (4) [Info](#) Refresh Actions ▼ Create route table

Find resources by attribute or tag < 1 > ⚙

<input type="checkbox"/>	Name ▼	Route table ID ▼	Explicit subnet associations	Edge associations	Main ▼	VPC ▼	Owner ID ▼
<input type="checkbox"/>	-	rtb-06a82450b99339ed8	-	-	Yes	vpc-0ee45b8b41c5f8e0b Main...	99627892...
<input type="checkbox"/>	Private Route Table	rtb-078092c209ea35632	subnet-02dc9a55261263444 / Subnet-2	-	No	vpc-0ee45b8b41c5f8e0b Main...	99627892...
<input type="checkbox"/>	Public Route Table	rtb-0a688f7202c9b2883	2 subnets	-	No	vpc-0ee45b8b41c5f8e0b Main...	99627892...

Security group

```
# Create security group for web traffic
resource "aws_security_group" "web" {
  name     = "WebSG"
  vpc_id   = aws_vpc.main.id

  # Ingress rules for ports 22, 80, 443, 8080
  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port = 443
    to_port   = 443
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port = 8080
    to_port   = 8080
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

A security group named "WebSG" has been created. This group permits incoming SSH, HTTP, HTTPS, and custom web traffic on port 8080 from any source. Since it's configured for a server that is not expected to initiate outbound traffic, it is set to allow all outgoing traffic by default

sg-07b306d920152b531 - WebSG Actions ▾

Details

Security group name
WebSG

Security group ID
sg-07b306d920152b531

Description
Managed by Terraform

VPC ID
vpc-0ee45b8b41c5f8e0b

Owner
996278926886

Inbound rules count
4 Permission entries

Outbound rules count
1 Permission entry

Inbound rules Outbound rules Tags

You can now check network connectivity with Reachability Analyzer Run Reachability Analyzer ×

Inbound rules (4) Manage tags Edit inbound rules

Filter security group rules

<input type="checkbox"/>	Name ▾	Security group rule... ▾	IP version ▾	Type ▾	Protocol ▾	Port range ▾	Source ▾	Description
<input type="checkbox"/>	-	sgr-0e00eb245c71e22...	IPv4	HTTP	TCP	80	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-0f5b5e398bffa74	IPv4	Custom TCP	TCP	8080	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-0bda670d8c7093d...	IPv4	SSH	TCP	22	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-053dd66d156460...	IPv4	HTTPS	TCP	443	0.0.0.0/0	-

Launch configuration

```
# Create Launch Configuration
resource "aws_launch_configuration" "my_config" {
  name          = "MyLaunchConfig"
  image_id      = "ami-024e6efaf93d85776"
  instance_type = "t2.micro"
  security_groups = [aws_security_group.web.id]

  user_data = <<-EOT
    #!/bin/bash
    sudo snap install docker
    sudo apt install git
    sudo docker pull projadt/weeb:final
    sudo docker run -d -p 8080:80 projadt/weeb:final
    sudo docker exec my_container sed -i "/Rohit Wagh/a <p>IP
Address: $(hostname -i)</p>" /usr/src/app/index.html
    sudo wget
https://raw.githubusercontent.com/finalprojadt/iadt/main/cpu.py
    sudo wget
https://raw.githubusercontent.com/finalprojadt/iadt/main/mykey.pub
    sudo cat /mykey.pub >> /home/ubuntu/.ssh/authorized_keys
  EOT

  lifecycle {
    create_before_destroy = true
  }
}
```

An EC2 instance is provisioned using the Ubuntu 22.04 image with a "t2.micro" specification and is associated with the "WEBSG" security group. In EC2 user data (boot strap script), we have installed docker and pulled an imager from docker hub and subsequently run, mapping its internal port 80 to

the host's port 8080. Fetched both the CPU utilization code and the local host's public key. This key was then added to the instance's authorized keys, allowing SSH access from the local host. Furthermore, we enabled the `create_before_destroy=true` setting to ensure immutability."

Instance summary for i-06406d25aa6a2aefa (MyAutoScalingGroup) [Info](#)

Refresh

Connect

Instance state ▼

Actions ▼

<div>Instance ID i-06406d25aa6a2aefa (MyAutoScalingGroup)</div> <div>IPv6 address -</div> <div>Hostname type IP name: ip-10-0-0-232.us-east-2.compute.internal</div> <div>Answer private resource DNS name -</div> <div>Auto-assigned IP address 3.22.66.145 [Public IP]</div> <div>IAM Role -</div> <div>IMDSv2 Optional</div>	<div>Public IPv4 address 3.22.66.145 open address</div> <div>Instance state Running</div> <div>Private IP DNS name (IPv4 only) ip-10-0-0-232.us-east-2.compute.internal</div> <div>Instance type t2.micro</div> <div>VPC ID vpc-0ee45b8b41c5f8e0b (Main VPC)</div> <div>Subnet ID subnet-0597c15362ddcb84e (Subnet-0)</div>	<div>Private IPv4 addresses 10.0.0.232</div> <div>Public IPv4 DNS ec2-3-22-66-145.us-east-2.compute.amazonaws.com open address</div> <div>Elastic IP addresses -</div> <div>AWS Compute Optimizer finding Opt-in to AWS Compute Optimizer for recommendations. Learn more</div> <div>Auto Scaling Group name terraform-2023081801070616610000002</div>
---	---	--

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

▼ Instance details [Info](#)

Platform
Ubuntu (Inferred)

Platform details
Linux/UNIX

Stop protection
Disabled

Instance auto-recovery
Default

AMI ID
ami-024e6efaf93d85776

AMI name
ubuntu/images/hvm-ssd/ubuntu-jammy-22.04-amd64-server-20230516

Launch time
Thu Aug 17 2023 21:07:10 GMT-0400 (Eastern Daylight Time) (about 17 hours)

Lifecycle
normal

Monitoring
detailed

Termination protection
Disabled

AMI location
amazon/ubuntu/images/hvm-ssd/ubuntu-jammy-22.04-amd64-server-20230516

Stop-hibernate behavior
disabled

Application Load Balancer

```
# ALB Security Group
resource "aws_security_group" "alb_sg" {
  name     = "ALB_SG"
  vpc_id = aws_vpc.main.id

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

resource "aws_lb" "my_alb" {
```

```

name                = "my-alb"
internal            = false
load_balancer_type  = "application"
security_groups     = [aws_security_group.alb_sg.id]
subnets            = [aws_subnet.subnets[0].id, aws_subnet.subnets[1].id]

enable_deletion_protection    = false
enable_cross_zone_load_balancing = true
}

resource "aws_lb_listener" "front_end" {
  load_balancer_arn = aws_lb.my_alb.arn
  port              = "80"
  protocol          = "HTTP"
  default_action {
    type = "forward"
    target_group_arn = aws_lb_target_group.my_tg.arn
  }
}

resource "aws_lb_target_group" "my_tg" {
  name      = "my-tg"
  port      = 8080
  protocol  = "HTTP"
  vpc_id    = aws_vpc.main.id
}

```

Another security group “aws_security_group”. This group is associated with Application Load Balancer(ALB) which only permits incoming traffic on port 80 and unrestricted outgoing traffic(Least Privilege principle) and configured to operate across two subnets. The ALB listens on port 80 for HTTP traffic and forwards it to a target group named 'my-tg', which expects traffic on port 8080.

sg-0026d89ee5ef6248c - ALB_SG Actions ▾

Details			
Security group name ALB_SG	Security group ID sg-0026d89ee5ef6248c	Description Managed by Terraform	VPC ID vpc-0ee45b8b41c5f8e0b ↗
Owner 996278926886	Inbound rules count 1 Permission entry	Outbound rules count 1 Permission entry	

Inbound rules | Outbound rules | Tags

📘 You can now check network connectivity with Reachability Analyzer Run Reachability Analyzer ✕

Inbound rules (1/1) 🔄 Manage tags Edit inbound rules

<input checked="" type="checkbox"/>	Name	Security group rule...	IP version	Type	Protocol	Port range	Source	Descri
<input checked="" type="checkbox"/>	-	sgr-Od1ff52eb2eeb0e8f	IPv4	HTTP	TCP	80	0.0.0.0/0	-

▼ Details

Load balancer type

Application

Status

Active

VPC

vpc-0ee45b8b41c5fbc0b

IP address type

IPv4

Scheme

Internet-facing

Hosted zone

Z3AADJGX6KTL2

Availability Zones

subnet-0597c15362ddcb84e us-east-2a (use2-az1)
subnet-0f7a149ea3434e50f us-east-2b (use2-az2)

Date created

August 17, 2023, 21:07 (UTC-04:00)

Load balancer ARN

arn:aws:elasticloadbalancing:us-east-2:996278926886:loadbalancer/app/my-alb/11433122a338b814

DNS name

my-alb-2121864173.us-east-2.elb.amazonaws.com (A Record)

Listeners and rules (1) Info

Manage rules

Manage listener

Add listener

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Filter listeners by property or value

< 1 > ⚙

Protocol:Port

▼

Default action

▼

Rules

▼

ARN

▼

Security policy

▼

Default SSL cert

▼

Tags

▼

HTTP:80

Forward to target group

• my-tg: 1 (100%)

• Group-level stickiness: Off

1 rule

ARN

Not applicable

Not applicable

0 tags

my-tg

Actions

Details

arn:aws:elasticloadbalancing:us-east-2:996278926886:targetgroup/my-tg/2122565a30f12bc6

Target type

Instance

Protocol : Port

HTTP: 8080

Protocol version

HTTP1

VPC

vpc-0ee45b8b41c5fbc0b

IP address type

IPv4

Load balancer

my-alb

Total targets

Healthy

Unhealthy

Unused

Initial

Draining

1

1

0

0

0

0

► Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

Targets

Monitoring

Health checks

Attributes

Tags

Registered targets (1)

Deregister

Register targets

Filter resources by property or value

< 1 > ⚙

Instance ID

▼

Name

▼

Port

▼

Zone

▼

Health status

▼

Health status details

i-06406d25aa6a2aefa

MyAutoScalingGroup

8080

us-east-2a

healthy

ROUTE 53

Route 53 > Hosted zones > projiadtl.online

Public **projiadtl.online** Info Delete zone Test record Configure query logging

► **Hosted zone details** Edit hosted zone

Records (3) DNSSEC signing Hosted zone tags (0)

Records (3) Info
Automatic mode is the current search behavior optimized for best filter results. [To change modes go to settings.](#)

↻ Delete record Import zone file Create record

Type ▼ Routing policy ▼ Alias ▼ < 1 > ⚙

Record ... ▼	Type ▼	Routin... ▼	Differe... ▼	Alias ▼	Value/Route traffic to ▼	TTL (s... ▼
projiadtl.o...	A	Simple	-	Yes	my-alb-2121864173.us-east-2.elb.amazonaws.com.	-
projiadtl.o...	NS	Simple	-	No	ns-1750.awsdns-26.co.uk. ns-245.awsdns-30.com. ns-840.awsdns-41.net. ns-1220.awsdns-24.org.	172800
projiadtl.o...	SOA	Simple	-	No	ns-1750.awsdns-26.co.uk. awsdns-hostmaster.ama...	900

< >

AUTO SCALING

```
# Create AutoScaling Group
resource "aws_autoscaling_group" "my_asg" {
  launch_configuration = aws_launch_configuration.my_config.name
  min_size             = 1
  max_size             = 5
  desired_capacity     = 1
  vpc_zone_identifier  = [aws_subnet.subnets[0].id, aws_subnet.subnets[1].id]
  target_group_arns    = [aws_lb_target_group.my_tg.arn]

  tag {
    key           = "Name"
    value         = "MyAutoScalingGroup"
    propagate_at_launch = true
  }
}

# Scale up policy
resource "aws_autoscaling_policy" "scale_up" {
  name              = "scale-up"
  scaling_adjustment = 1
  adjustment_type   = "ChangeInCapacity"
```

```

    cooldown                = 180
    autoscaling_group_name = aws_autoscaling_group.my_asg.name
}

# Scale down policy
resource "aws_autoscaling_policy" "scale_down" {
    name                = "scale-down"
    scaling_adjustment  = -1
    adjustment_type     = "ChangeInCapacity"
    cooldown            = 180
    autoscaling_group_name = aws_autoscaling_group.my_asg.name
}

# CloudWatch Alarm to scale up
resource "aws_cloudwatch_metric_alarm" "scale_up_alarm" {
    alarm_name                = "scale-up-alarm"
    comparison_operator       = "GreaterThanOrEqualToThreshold"
    evaluation_periods        = "1"
    metric_name               = "CPUUtilization"
    namespace                 = "AWS/EC2"
    period                    = "30"
    statistic                 = "Average"
    threshold                 = "70"
    alarm_description         = "This metric triggers when CPU usage exceeds 70%"
    alarm_actions              = [aws_autoscaling_policy.scale_up.arn]
}

# CloudWatch Alarm to scale down
resource "aws_cloudwatch_metric_alarm" "scale_down_alarm" {
    alarm_name                = "scale-down-alarm"
    comparison_operator       = "LessThanOrEqualToThreshold"
    evaluation_periods        = "1"
    metric_name               = "CPUUtilization"
    namespace                 = "AWS/EC2"
    period                    = "30"
    statistic                 = "Average"
    threshold                 = "40"
    alarm_description         = "This metric triggers when CPU usage falls below 40%"
    alarm_actions              = [aws_autoscaling_policy.scale_down.arn]
}

```

Established an AWS Auto Scaling Group “my_asg” that dynamically adjusts instance counts between 1 and 5 based on CPU usage. It monitor the CPU usage of all the instance that is running and If CPU exceeds 70%(average), it scales up by one instance, and if it drops below 40%(average), it scales down by one. These decisions are monitored and triggered by CloudWatch Alarms.

terraform-20230818010706166100000002

[Details](#) | [Activity](#) | [Automatic scaling](#) | [Instance management](#) | [Monitoring](#) | [Instance refresh](#)

Group details

Edit

Auto Scaling group name terraform-20230818010706166100000002	Desired capacity 1	Status -	Amazon Resource Name (ARN) arn:aws:autoscaling:us-east-2:996278926886:au toScalingGroup:03b6a4e6-47d9-4329-9f37-cd42ba b9fbe8:autoScalingGroupName/terraform-2023081 8010706166100000002
Date created Thu Aug 17 2023 21:07:06 GMT-0400 (Eastern Daylight Time)	Minimum capacity 1		
	Maximum capacity 5		

Launch configuration

Edit

Launch configuration MyLaunchConfig	AMI ID ami-024e6efaf93d85776	Instance type t2.micro	Create time Thu Aug 17 2023 21:06:57 GMT-0400 (Eastern Daylight Time)
Storage (volumes) -	Security groups sg-07b306d920152b531	Key pair name -	
View details in the launch configuration console			

Network

Edit

Availability Zones us-east-2a, us-east-2b	Subnet ID subnet-0597c15362ddcb84e, subnet- 0f7a149ea3434e50f	
--	---	--

scale-down



Simple scaling

Enabled

scale-down-alarm

breaches the alarm threshold: CPUUtilization =< 40 for 1 consecutive periods of 30 seconds for the metric dimensions:

Remove 1 capacity units

180 seconds before allowing another scaling activity

scale-up



Simple scaling

Enabled

scale-up-alarm

breaches the alarm threshold: CPUUtilization >= 70 for 1 consecutive periods of 30 seconds for the metric dimensions:

Add 1 capacity units

180 seconds before allowing another scaling activity

Alarms (2)



Any state

Any type

Any actions status

☐ Hide Auto Scaling alarms

< 1 >

scale-down-alarm

Metric alarm

Insufficient data

scale-up-alarm

Metric alarm


Insufficient data

scale-up-alarm



View

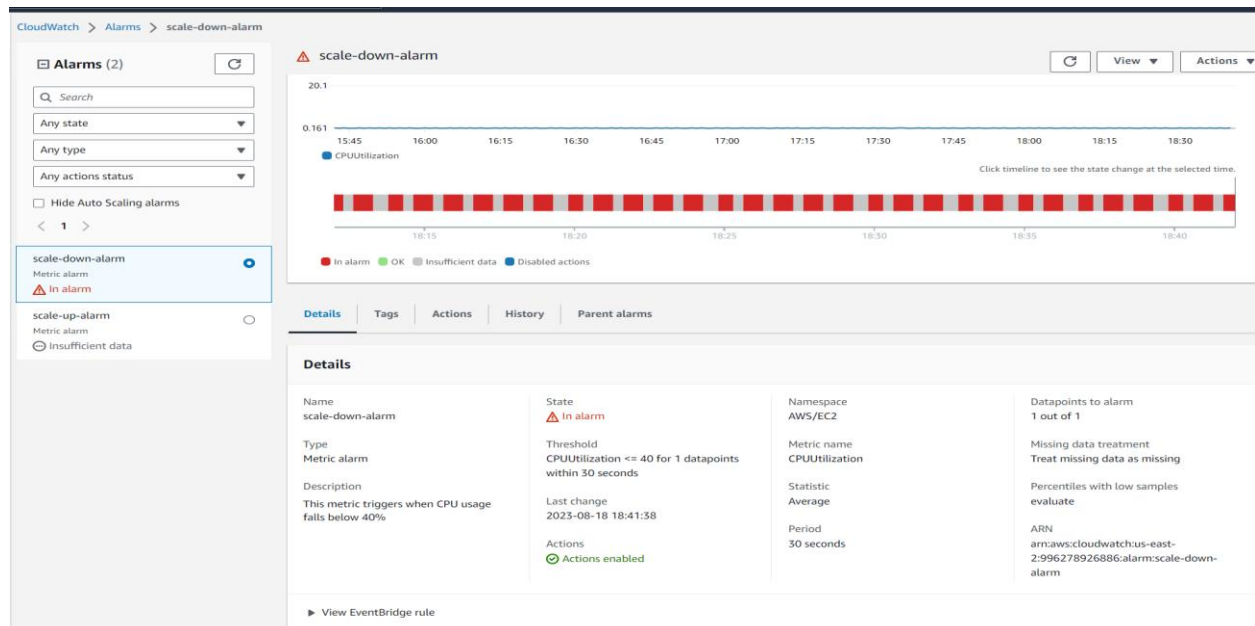
Actions



[Details](#) | [Tags](#) | [Actions](#) | [History](#) | [Parent alarms](#)

Details

Name scale-up-alarm	State Insufficient data	Namespace AWS/EC2	Datapoints to alarm 1 out of 1
Type Metric alarm	Threshold CPUUtilization >= 70 for 1 datapoints within 30 seconds	Metric name CPUUtilization	Missing data treatment Treat missing data as missing
Description This metric triggers when CPU usage exceeds 70%	Last change 2023-08-18 18:41:17	Statistic Average	Percentiles with low samples evaluate
	Actions Actions enabled	Period 30 seconds	ARN arn:aws:cloudwatch:us-east- 2:996278926886:alarms:scale-up-alarm



CPU LOAD

```
import multiprocessing
import time

def load_cpu(load_percentage, interval):
    """
    A function that generates a specific load on the CPU.
    """
    work_time = interval * load_percentage
    sleep_time = interval - work_time

    while True:
        end_time = time.time() + work_time
        while time.time() < end_time:
            x = (0.00001*3.14*3.14) / 2.34
            time.sleep(sleep_time)

NUM_PROCESSES = 1
LOAD_PERCENTAGE = 0.80 # 80% load
INTERVAL = 1.0 # interval in seconds

if __name__ == "__main__":
    processes = []

    for _ in range(NUM_PROCESSES):
```

```

        process = multiprocessing.Process(target=load_cpu, args=(LOAD_PERCENTAGE,
INTERVAL))
        processes.append(process)
        process.start()

try:
    while True:
        time.sleep(1) # Keep the script running
except KeyboardInterrupt:
    for process in processes:
        process.terminate()


```

This code generates CPU load by running a computational task in parallel processes. It uses the multiprocessing module to spawn number of processes, each applying an 80% load on the CPU over 1-second intervals. This is achieved by performing calculations for 80% of the time and resting for the remaining 20%. The main script keeps running indefinitely, and upon manual interruption, all active processes are terminated.

Web Application

Developed a Tic Tac Toe game using HTML and then packaged it into a container. This containerized version was then uploaded to Docker Hub

(<https://hub.docker.com/repository/docker/projiadt/weeb/general>)



The screenshot shows the Docker Hub interface for the repository 'projiadt / weeb'. The page includes a search bar, navigation links, and a 'General' tab. A description field is present with a placeholder text. The 'Tags' section lists three tags: 'final', 'lat', and 'latest', each with its OS, type, and push/pull times. The 'Automated Builds' section provides information on connecting to GitHub or Bitbucket for automated builds.

projiadt / weeb

Description
This repository does not have a description

Tags
This repository contains 3 tag(s).

Tag	OS	Type	Pulled	Pushed
final	linux	Image	19 hours ago	3 days ago
lat	linux	Image	3 days ago	4 days ago
latest	linux	Image	4 days ago	4 days ago

Automated Builds
Manually pushing images to Hub? Connect your account to GitHub or Bitbucket to automatically build and tag new images whenever your code is updated, so you can focus your time on creating.

Available with Pro, Team and Business subscriptions. [Read more about automated builds](#)

Tic Tac Toe

X		O
X	O	X
O	X	O

O wins!


Reset Game

Project by Shah Nawaz Syed Shah, Ali Goudarzi, and Rohit Wagh

Code link : <https://github.com/finalprojiadt/iadt/blob/main/dummy.html>

Domain Name:

Domains → Details

projiadt.online

Domain

Products

Sharing & Transfer

Advanced DNS

STATUS & VALIDITY


?

✓ ACTIVE

Aug 15, 2023 - Aug 15, 2024

☐ AUTO-RENEW

ADD YEARS

WithheldforPrivacy

?


☒ PROTECTION

Aug 15, 2023 - Aug 15, 2024

☐ AUTO-RENEW

ADD YEARS

SHOW DETAILS

PremiumDNS

?

Enable PremiumDNS protection in order to switch your domain to our PremiumDNS platform. With our PremiumDNS platform, you get 100% DNS uptime and DDoS protection at the DNS level.

BUY NOW

NAMESERVERS

?

Custom DNS

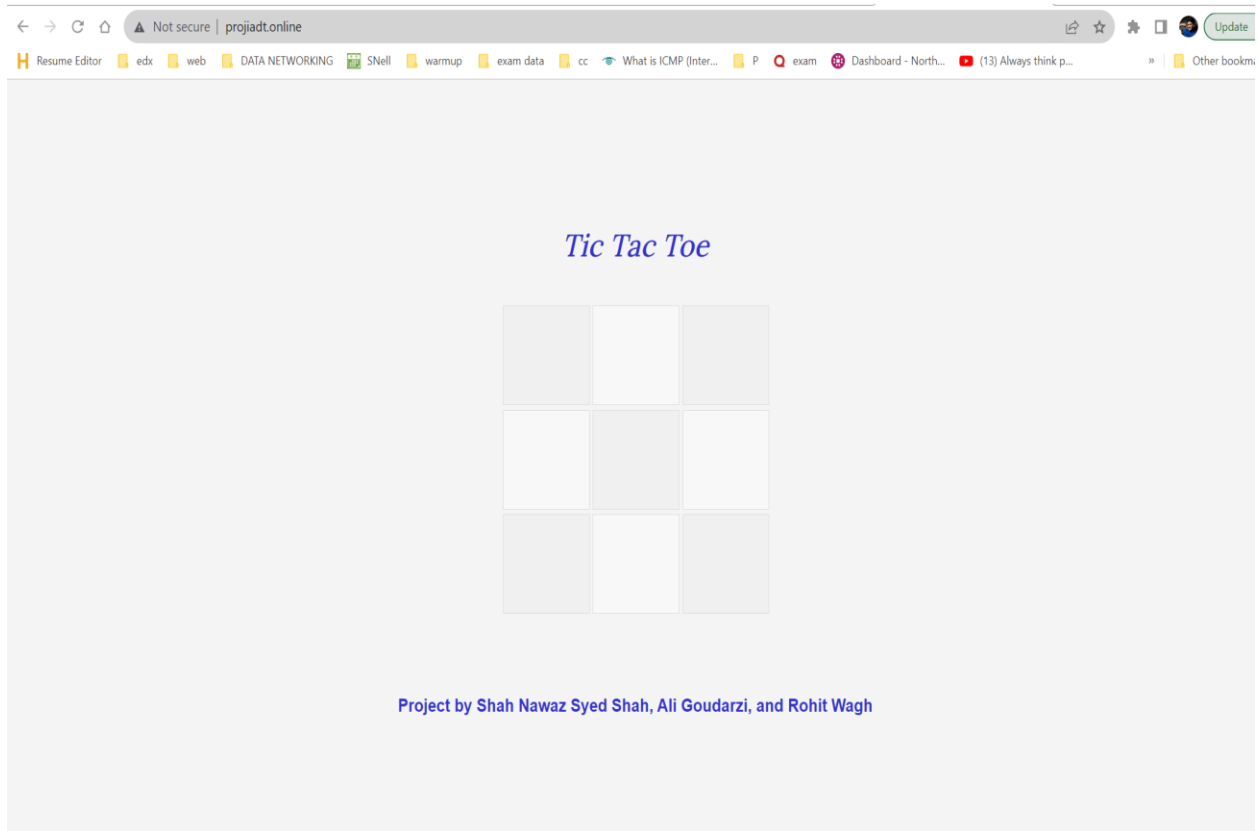
ns-527.awsdns-01.net.

ns-1621.awsdns-10.co.uk.

ns-488.awsdns-61.com.

ns-1161.awsdns-17.org.

ADD NAMESERVER



DIFFICULTIES:

1. SSH into each instance to increase the CPU load

Though we developed the script to increase the CPU load and tested it, for each ec2 instance which is being automatically created using Auto Scaling Group, we were not able to SSH into them.

In order to Overcome this problem, we generated a SSH key and pushed that public key into each instance using EC2 boot strap script. After accomplishing this, we were able to increase the CPU load.

2. Auto Scaling

Initially we had the auto scaling to scale up when any one of the CPU increases above the 70%. But, on a longer run we realized that load will be distributed and average CPU load has to be calculated. So later we changed, the measuring criteria from Maximum to Average.

3. Route 53

Initially, we employed Terraform for DNS automation. However, the challenge we encountered was that every time we executed terraform apply, a new elastic load balancer and Route53 instance were generated. As a result, we consistently had to update the settings on namecheap.com, leading to a 24-48 hour delay for projadt.online to become active.

BONUS:

1. Immutability

Immutability means not updating/ modifying the resource once they have been deployed.

Rather replacing the older one with a new one. This prevents configuration drift

In our project by using the

“create_before_destroy=true”, we ensure that new version of the instance is created before the old version is destroyed and this ensures that there is no downtime.

2. Elastic Load Balancer

We employed an Application Load Balancer to evenly distribute incoming traffic across our web servers. As it's internet-facing, it guarantees high availability, ensuring the application remains accessible even if one of the instances fails.

3. Container

We've encapsulated our web application within a container, streamlining its deployment.

Future Scope:

1. RDS & NAT instance/ NAT Gateway

We can utilize RDS to record the game's winners and also keep track of IP address of each container and place the RDS within a private subnet. This ensures that only instances within our network can interact with it. If we need to provide internet access to this RDS, we can do so using a NAT instance (using Bastion Host) or a NAT Gateway.

2. NACL

Though we have security group which prevents malicious activity, being stateful, if server presses by mistake it will allow the traffic in without checking. By having NACL and its stateful as well, network will have additional security.

3. Custom AMI

Using the current Terraform script, every deployment involves setting up software packages and other resources, leading to longer boot-up times for the EC2 instances. If we instead use a customized AMI with Docker and other required software pre-installed, the instances would launch faster. The more commands and data you have in EC2's user data, the longer it takes for the instance to become operational. Leveraging a pre-configured AMI can optimize this process.

4. Ansible Dynamic Inventory

Instead of executing scripts to manually increase the CPU load on each instance, we can utilize Dynamic Inventory. This approach allows for more efficient provisioning and monitoring of each EC2 instance.